# **DV2575: Projects**

Yong Yao

yong.yao@bth.se

DIDD, BTH



#### General

- Introduction
  - CUDA programming
- Project 1
  - Odd-even sort
- □ Project 2
  - Gaussian elimination
- □ Project 3
  - Ray tracing



- CUDA
  - Compute Unified Device Architecture
  - Freely distributed by NVIDIA
- Basic idea
  - The GPU is linked to the CPU by a reasonably fast connection
  - Use the GPU as a co-processor
- Further
  - Farm out big parallel tasks to the GPU
  - Keep the CPU busy with the control of the execution and "corner" tasks

- CUDA installation
  - Support various Operating Systems (Oss)
  - I.e., Windows, Mac OSX, Linux
- Quick start guide
  - See document "CUDA\_Quick\_Start\_Guide\_v7.5.pdf" on the course page in its-learning
- More questions
  - □ Feel free to ask Yong Yao (in person) in Lab or office H453A
  - Note that, will not answer by e-mail



- CUDA programming: extended C
- Declaration specifications:
  - global, device, shared, local, constant
- Keywords
  - threadldx, blockldx
- □ Intrinsics
  - syncthreads
- Runtime API
  - For memory and execution management
- Kernel launch

```
device float filter[N];
  global void convolve (float
*image) {
  shared float region[M];
  region[threadIdx.x] = image[i];
  syncthreads()
  image[j] = result;
// Allocate GPU memory
void *myimage = cudaMalloc(bytes)
// 100 blocks, 10 threads per
block
convolve<<<100, 10>>> (myimage);
```

Example

```
__global__ void testKernel (void) {
int main(void) {
  testKernel<<<1,1>>>();
  return 0;
}
```

- global\_\_\_
  - CUDA C/C++ keyword
  - The function runs on the device, and is called from host code
- main()
  - Processed by standard host compiler, e.g., gcc
- testKernel()
  - Processed by NVIDIA compiler



- □ testKernel<<<1,1>>>();
  - Triple angle brackets mark a call from host code to device code
  - This is required to execute a function on the GPU
- <<<1,1>>>
  - The host instructs an execution configuration
  - At GPU, it is supposed to run 1 block, each with 1 thread
  - A block represents the entity that gets executed by an SM (stream multiprocessor)
  - Each block cannot has larger than 1024 threads
- □ <<<x,y,z>>>
  - 3D structure



- CUDA function declarations
  - global: must return void
  - Others, see CUDA reference manual

	Executed on the	Only callable from
device float myDeviceFunc()	device	device
global void myKernelFunc()	device	host
host float myHostFunc()	host	host



- Important APIs
- cudaMalloc()
  - Allocates object in the device Global Memory
  - Para1: address of a pointer to the allocated object
  - Para2: size of allocated object
- cudaFree()
  - Frees object from device Global Memory
  - Para1: pointer to freed object
- cudaMemcpy()
  - Frees object from device Global Memory
  - Pointer to freed object



- Important APIs (cont.)
- cudaMemcpy()
  - Memory data transfer
  - Para 1: pointer to source
  - Para 2: pointer to destination
  - Para 3: number of bytes copied
  - Para 4: type of transfer
    - Host to Host
    - Host to Device
    - Device to Host
    - Device to Device



#### Matrix manipulation

```
global void MatrixMulKernel(Matrix M, Matrix N,
Matrix P) {
 int tx = threadIdx.x;
 int ty = threadIdx.y;
 float Pvalue = 0;
 for (int k = 0; k < M.width; ++k) {
                                                                         N
   float Melement = M.elements[ty * M.width + k];
   float Nelement = N.elements[k * N. width + tx];
   Pvalue += Melement * Nelement;
 P.elements[ty * P. width + tx] = Pvalue;
}
                                                      M
                                                                         P
                                                                                    tx
void main() {
 MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd);
                                                                            ty
                                                            Width
```

#### Project 1: odd-even sort

#### General concept

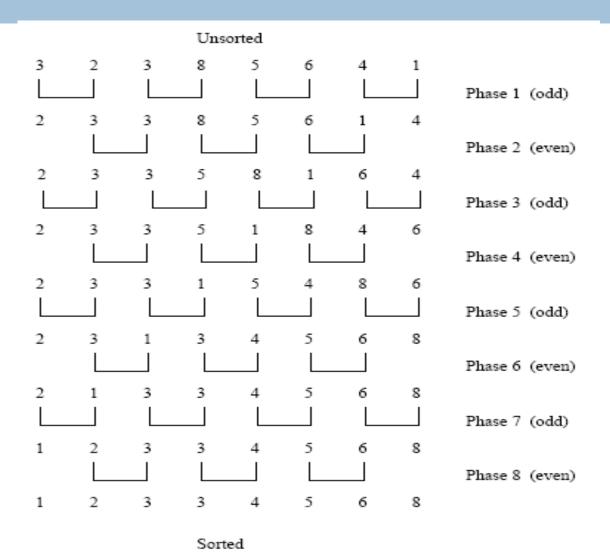
- An odd-even sort is a comparison sort related to bubble sort.
- It functions by comparing all odd/even-indexed pairs of adjacent elements in the unsorted list of numbers.

#### Comparison goal

- Make each pair of numbers in the expected order by doing switch operation
- (e.g., the first number is smaller than the second one) n

#### Comparison rules

- Repeating the comparison for each even(/odd)-indexed pairs
- Alternating the comparison between odd->even and even->odd steps.



## Project 1: odd-even transposition

- □ The algorithm is guaranteed to terminate after N/2 odd-even and even-odd steps.
- After n phases of odd-even exchanges, the sequence is sorted.
- fineq Each phase of the algorithm (either odd or even) requires  $\Theta(n)$  comparisons.
- $\square$  Serial complexity is  $\Theta(n^2)$ .



### Project 1: parallel transposition

- Consider the one item per processor case.
- There are n iterations, in each iteration, each processor does one compare-exchange.
- luleq The parallel run time of this formulation is  $\Theta\left(n
  ight)$ .
- This is cost optimal with respect to the base serial algorithm but not the optimal one.



### Project 1: tasks

- Implement the single thread based odd-even sort
- Implement the multiple-thread based odd-even sort
- Compare the time used for soring in the two methods
- Performance analysis showing how the following factors affect the performance:
  - Effect of number of integers on sorting.
  - Effect of number of threads on sorting.
- One student per implementation
- Evaluation: G/U



### Project 2: gaussian elimination

#### General concept

- An algorithm to solve linear equations.
- Conduct a sequence of operations on the matrix with coefficients

#### Solution approach

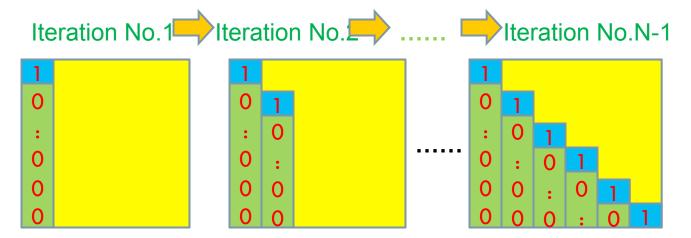
- Transform linear equations into an upper-triangular matrix
- A pivot column is used to modify the matrix until the lower lefthand corner is filled with zeros, as much as possible.
- After the transformation, back-substitution is applied.

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1k} \\ a_{21} & a_{22} & \cdots & a_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ a_{k1} & a_{k2} & \cdots & a_{kk} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_k \end{bmatrix},$$



# Project 2: gaussian elimination

- Solution approach (cont.)
  - Forward substitution



**D** Back substitution, for vector  $(x_1, x_2, ..., x_{k-1}, x_k)$  then compute the result from  $x_k$  to  $x_1$ , i.e.,  $x_k -> x_{k-1} -> ... -> x_2 -> x_1$ 

$$\begin{cases} 3x - 2y + 2z = 16 \\ 7x - 3y + 2z = 26 \\ 2x - y + 4z = 18 \end{cases}$$



$$\begin{cases} 3x - 2y + 2z = 16 \\ 7x - 3y + 2z = 26 \end{cases} *(-2) + 2x - y + 4z = 18$$



$$\begin{cases} 3x - 2y + 2z = 16 \\ x + y - 2z = -6 \\ 2x - y + 4z = 18 \end{cases}$$



$$\begin{cases} x + y - 2z = -6 \\ 2x - y + 4z = 18 \\ 3x - 2y + 2z = 16 \end{cases}$$



$$\begin{cases} x + y - 2z = -6 \\ 2x - y + 4z = 18 \\ 3x - 2y + 2z = 16 \end{cases} *(-2) +$$



$$\begin{cases} x + y - 2z = -6 \\ -3y + 8z = 30 \\ -5y + 8z = 34 \end{cases} * \left(-\frac{5}{3}\right) +$$



$$\begin{cases} x + y - 2z = -6 \\ -3y + 8z = 30 \end{cases}$$
$$-\frac{16}{3}z = -16$$




#### **Back substitution**

$$\begin{cases} x + y - 2z = -6 \\ -3y + 8z = 30 \end{cases} \begin{cases} z = 3 \\ y = -2 \\ x = 2 \end{cases}$$



## Project 2: CUDA based solution


□ Hits: combine forward- and back-substitution





#### Project 2: CUDA based solution

□ Hits: combine forward- and back-substitution



### Project 2: tasks

- Implement the single thread based gaussian elimination
- Implement the multiple-thread based gaussian elimination
- Compare the time used in the two methods
- Performance analysis showing how the following factors affect the performance:
  - Effect of number of variables on solving linear equations.
  - Effect of number of threads on solving linear equations.
- One or two students per implementation
- Evaluation: G/U



## Project 3: ray tracing

#### Ray tracing

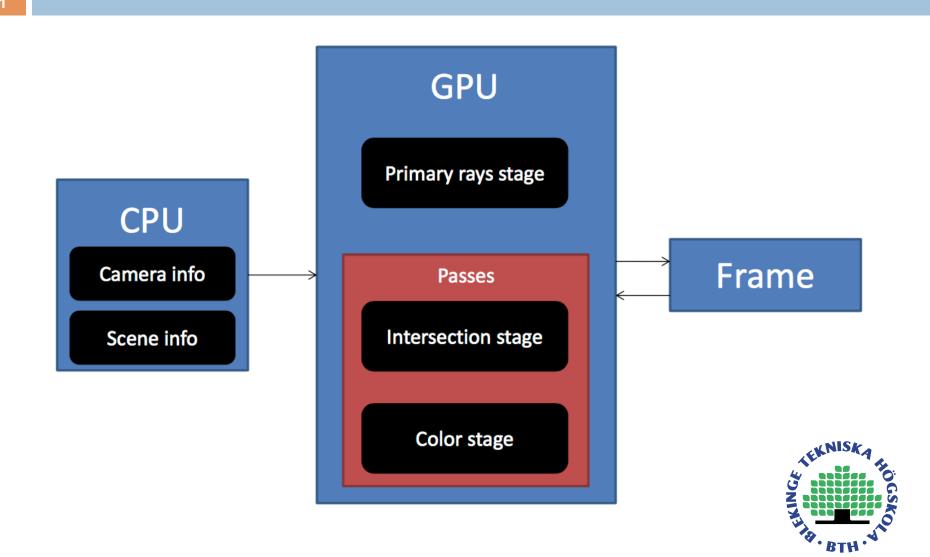
- It is often used to render realistic images where rendering time is less important.
- By using current generation of GPGPU hardware it is possible to perform ray tracing in real time.

#### Inspired by book chapters

- Interactive Ray Tracing Using the Compute Shader in DirectX 11
- Bit-Trail Traversal for Stackless LBVH on DirectCompute



# Project 3: system



### Project 3: tasks

- Generate primary rays from a camera position and orientation.
- Support diffuse lighting with light attenuation Implement the single thread based odd-even sort
- Support multiple ray traces/bounces Implement the multiple-thread based odd-even sort
- Support ray tracing of one triangle mesh



## Project 3: tasks (cont.)

- Performance analysis showing how the following factors affect the performance:
  - Number of threads per thread group.
  - Screen resolution.
  - Trace depth.
  - Number of light sources.
- One or two students per implementation
- □ Evaluation: G/U



## Project 3: Example 1

10 lights - 1 trace
 A good tip is to generate the box automatically (same buffer)!





# Project 3: Example 2

□ 10 lights - 10 trace

